

MV-ADAPTER: MULTI-VIEW CONSISTENT IMAGE GENERATION MADE EASY

Zehuan Huang¹, Yuan-Chen Guo^{2†}, Haoran Wang³, Ran Yi³, Lizhuang Ma³, Yan-Pei Cao^{2✉}, Lu Sheng^{1✉}

¹School of Software, Beihang University ²VAST ³Shanghai Jiao Tong University
Project page: <https://huanngzh.github.io/MV-Adapter-Page/>

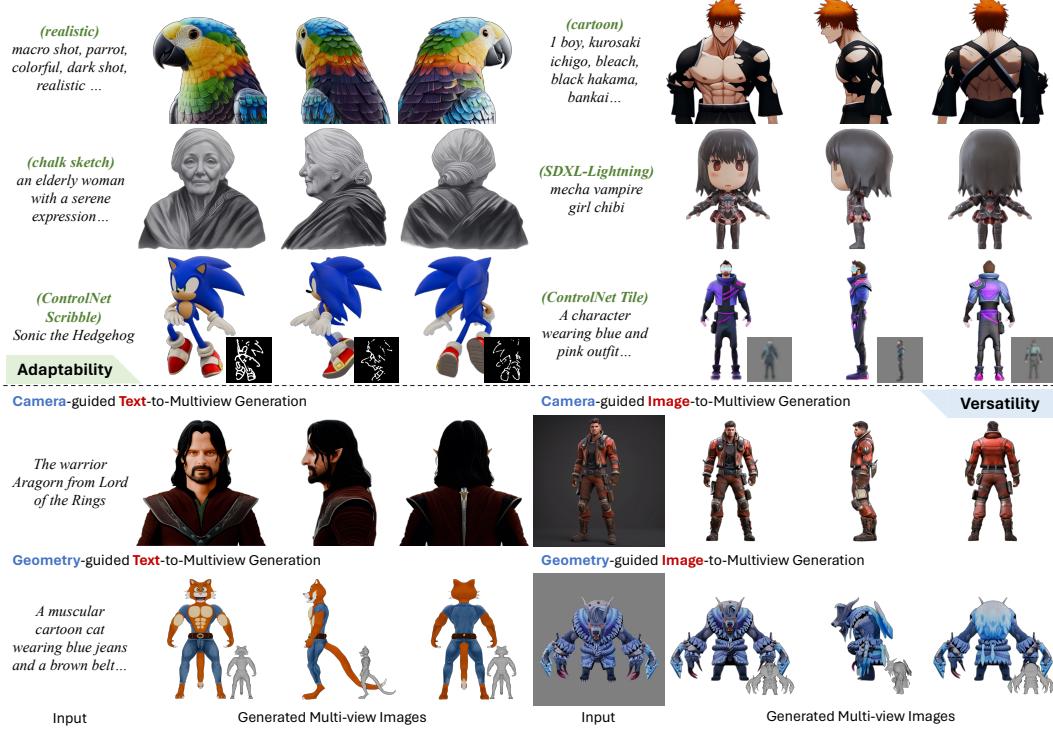


Figure 1: MV-Adapter is a versatile plug-and-play adapter that turns existing pre-trained text-to-image (T2I) diffusion models to multi-view image generators. **Row 1,2,3:** results by integrating MV-Adapter with personalized T2I models, distilled few-step T2I models, and ControlNets (Zhang et al., 2023), demonstrating its **adaptability**. **Row 4,5:** results under various control signals, including view-guided or geometry-guided generation with text or image inputs, showcasing its **versatility**.

ABSTRACT

Existing multi-view image generation methods often make invasive modifications to pre-trained text-to-image (T2I) models and require full fine-tuning, leading to (1) high computational costs, especially with large base models and high-resolution images, and (2) degradation in image quality due to optimization difficulties and scarce high-quality 3D data. In this paper, we propose the **first adapter-based solution for multi-view image generation**, and introduce MV-Adapter, a versatile plug-and-play adapter that enhances T2I models and their derivatives **without altering the original network structure or feature space**. By updating fewer parameters, MV-Adapter enables efficient training and preserves the prior knowledge embedded in pre-trained models, mitigating overfitting risks. To efficiently model the 3D geometric knowledge within the adapter, we **introduce innovative designs that include duplicated self-attention layers and parallel attention architec-**

[✉]Corresponding authors: Isheng@buaa.edu.cn, caoyanpei@gmail.com; [†]: project lead

ture, enabling the adapter to inherit the powerful priors of the pre-trained models to model the novel 3D knowledge. Moreover, we present a unified condition encoder that seamlessly integrates camera parameters and geometric information, facilitating applications such as text- and image-based 3D generation and texturing. MV-Adapter achieves multi-view generation at 768 resolution on Stable Diffusion XL (SDXL), and demonstrates adaptability and versatility. It can also be extended to arbitrary view generation, enabling broader applications. We demonstrate that MV-Adapter sets a new quality standard for multi-view image generation, and opens up new possibilities due to its efficiency, adaptability and versatility.

1 INTRODUCTION

Multi-view image generation is a fundamental task with significant applications in areas such as 2D/3D content creation, robotics perception, and simulation. With the advent of text-to-image (T2I) diffusion models (Ramesh et al., 2022; Nichol et al., 2022; Saharia et al., 2022; Ramesh et al., 2021; Balaji et al., 2022; Podell et al., 2024; Mokady et al., 2023), there has been considerable progress in generating high-quality single-view images. Extending these models to handle multi-view generation holds the promise of unifying text, image, and 3D data into a cohesive framework.

Recent attempts on multi-view image generation (Shi et al., 2023b; Tang et al., 2023; 2024; Huang et al., 2024; Gao et al., 2024; Liu et al., 2023a; Long et al., 2024; Li et al., 2024; Kant et al., 2024; Zheng & Vedaldi, 2024; Wang & Shi, 2023) involve fine-tuning T2I models on large-scale 3D datasets (Deitke et al., 2023; Yu et al., 2023) and propose modeling 3D consistency across images by applying attention on relevant pixels in different views. However, this is computationally challenging when working with large base T2I models and high-resolution images, as it requires at least n view images to be processed simultaneously during training. Existing advanced methods (Li et al., 2023a; 2024) still struggle with 512 resolution, which is far from the 1024 or higher that modern T2I models can achieve. Moreover, the scarcity of high-quality 3D training data exacerbates the optimization difficulty when performing full model fine-tuning, resulting in a degradation in the quality of the generated multi-view images. These limitations primarily stem from the invasive changes to base models and full tuning.

To address these challenges, we propose the first adapter-based solution for multi-view image generation. The adapter mechanism plays a crucial role in this context for several reasons: First, adapters are easy to train. They require updating only a small number of parameters, making the training process faster and more memory-efficient. This property has become increasingly critical as state-of-the-art T2I models grow in scale, making full fine-tuning infeasible. Second, adapters help in preserving prior knowledge embedded in the pre-trained models. Adapters mitigate the risk of overfitting by constraining the optimization space through fewer trainable parameters, allowing the model to retain its learned priors while adapting to multi-view generation. Third, adapters offer adaptability and ease of use. They are plug-and-play modules and can be applied to different variants of base models, including fine-tuned versions (Ruiz et al., 2023) and LoRAs (Hu et al., 2021).

Building on the importance of adapters for the multi-view generation task and adhering to the principle of preserving the original network structure and feature space of the base T2I model, we propose MV-Adapter, a versatile plug-and-play adapter that enhances T2I models and their derivatives for multi-view generation under various conditions. To achieve this, we design an effective adapter framework with innovative features. Unlike existing methods (Shi et al., 2023b;a) that modify the base model’s self-attention layers to include multi-view or reference features, which disrupts learned priors and requires full model fine-tuning, we duplicate the self-attention layers to create new multi-view attention and image cross-attention layers, and initializes the output projections to zero. We further enhance the effectiveness of our attention layers through a parallel organization structure, ensuring that the new layers fully inherit the powerful priors of the pre-trained self-attention layers, thus enabling efficient learning of geometric knowledge. Additionally, we introduce a unified condition embedding and encoder that seamlessly integrates camera parameters and geometric information into spatial map representations, enhancing the model’s versatility and applicability.

By leveraging our adapter design, we successfully achieve the multi-view generation at 768 resolution on Stable Diffusion XL (SDXL) (Podell et al., 2024). As shown in Fig. 1, our trained MV-Adapter demonstrates both adaptability and versatility. It seamlessly applies to derivatives of the

base model (Ruiz et al., 2023; Hu et al., 2021; Zhang et al., 2023; Mou et al., 2024) for customized or controllable multi-view generation, while simultaneously supporting camera and geometry guidance, which benefits applications in 3D generation and texture generation. Moreover, MV-Adapter can be extended to arbitrary view generation, enabling broader applications.

In summary, our contributions are as follows: (1) We propose the first adapter-based approach that enhances efficiency and is able to work with larger base models for higher performance. (2) We introduce an innovative adapter framework that efficiently models 3D geometric knowledge and supports versatile applications like 3D generation and texture generation. (3) Our MV-Adapter can be extended to generate images from arbitrary viewpoints, facilitating a wider range of downstream tasks. (4) MV-Adapter provides a framework for decoupled learning that offers insights into modeling new types of knowledge, such as physical or temporal knowledge.

2 RELATED WORK

Text-to-image diffusion models. Text-to-image (T2I) generation (Ramesh et al., 2022; Nichol et al., 2022; Saharia et al., 2022; Ramesh et al., 2021; Balaji et al., 2022; Podell et al., 2024; Mokady et al., 2023) has made remarkable progress, particularly with the advancement of diffusion models (Ho et al., 2020; Song et al., 2020; Dhariwal & Nichol, 2021; Ho & Salimans, 2022). Guided diffusion (Dhariwal & Nichol, 2021) and classifier-free guidance (Ho & Salimans, 2022) improved text conditioning and generation fidelity. DALL-E2 (Ramesh et al., 2022) leverages CLIP (Radford et al., 2021) for better text-image alignment. The Latent Diffusion Model (Rombach et al., 2022), also known as Stable Diffusion, enhances efficiency by performing diffusion in the latent space of an autoencoder. Stable Diffusion XL (Podell et al., 2024), a two-stage cascade diffusion model, has greatly improved the generation of high-frequency details and image quality.

Derivatives and extensions of T2I models. To facilitate creation with pre-trained T2Is, various derivative models and extensions have been developed, focusing on model distillation for efficiency (Meng et al., 2023; Song et al., 2023; Luo et al., 2023; Lin et al., 2024) and controllable generation (Cao et al., 2024). These derivatives encompass personalization (Ruiz et al., 2023; Gal et al., 2022; Hu et al., 2021; Shi et al., 2024; Wang et al., 2024a; Ma et al., 2024; Song et al., 2024; Kumari et al., 2023; Ye et al., 2023), and spatial control (Mou et al., 2024; Zhang et al., 2023). Typically, they employ adapters or fine-tuning methods to extend functionality while preserving the original feature space of the pre-trained models. Our work adheres to non-intrusive principle, ensuring compatibility with these derivatives or extensions for broader applications.

Multi-view Generation with T2I models. Multi-view generation methods (Shi et al., 2023b; Tang et al., 2023; 2024; Huang et al., 2024; Gao et al., 2024; Liu et al., 2023a; Long et al., 2024; Li et al., 2024; Kant et al., 2024; Zheng & Vedaldi, 2024; Wang & Shi, 2023; Jeong et al., 2025) extend T2I models by leveraging large-scale 3D datasets (Deitke et al., 2023; Yu et al., 2023). For instance, MVDream (Shi et al., 2023b) integrates camera embeddings and expands the self-attention mechanism from 2D to 3D for cross-view connections, while SPAD (Kant et al., 2024) enhances spatial relational modeling by applying epipolar constraints to cross-view attention. Era3D (Li et al., 2024) introduces an efficient row-wise self-attention mechanism aligned with epipolar lines across views, facilitating high-resolution multi-view generation. However, these methods typically require extensive parameter updates, altering the feature space of pre-trained T2I models and limiting their compatibility with T2I derivatives. Our work addresses this by introducing a multi-view adapter that harmonizes with pre-trained T2Is, significantly expanding the potential for diverse applications.

3 PRELIMINARY

Here we introduce the preliminary of multi-view diffusion models (Shi et al., 2023b; Kant et al., 2024; Li et al., 2024), which can help understand the common strategies in modeling multi-view consistency within T2I models.

Multi-view diffusion models. Multi-view diffusion models enhance T2Is by introducing multi-view attention mechanism, enabling the generation of images that are consistent across different viewpoints. Several studies (Shi et al., 2023b; Wang & Shi, 2023) extend the self-attention of T2Is to include all pixels across multi-view images. Let f^{in} denotes the input of the attention block, the

dense multi-view self-attention extends f^{in} from the view itself to the concatenated feature sequence from n views. While this approach captures global dependencies, it is computationally intensive, as it processes all pixels of all views. To mitigate the computational cost, epipolar attention (Kant et al., 2024; Huang et al., 2024) leverages geometric relationships between views. Specifically, methods like SPAD (Kant et al., 2024) extend the self-attention by restricting f^{in} to the view itself as well as patches along its epipolar lines.

Furthermore, when generating orthographic views at an elevation angle of 0° , the epipolar lines align with the image rows. Utilizing this property, row-wise self-attention (Li et al., 2024) is introduced after the original self-attention layers in T2I models. The process is defined as:

$$f^{self} = \text{SelfAttn}(f^{in}) + f^{in}; f^{mv} = \text{MultiViewAttn}(f^{self}) + f^{self} \quad (1)$$

where MultiViewAttn performs attention across the same rows in different views, effectively enforcing multi-view consistency with reduced computational overhead.

4 METHOD

MV-Adapter is a plug-and-play adapter that learns multi-view priors transferable to derivatives of T2Is without specific tuning, and enable them to generate multi-view consistent images under various conditions. As shown in Fig. 2, at inference, our MV-Adapter, which contains a condition guider and the decoupled attention layers, can be inserted into a personalized or distilled T2I to constitute the multi-view generator.

In detail, as shown in Fig. 3, the condition guider in Sec. 4.1 encodes the camera or geometry information, which supports both camera-guided and geometry-guided generation. Within the decoupled attention mechanism in Sec. 4.2, the additional multi-view attention layers learn multi-view consistency, while the optional image cross-attention layers are for image-conditioned generation. These new layers are duplicated from pre-trained spatial self-attention and organized in a parallel architecture. Sec. 4.3 elaborates on the training and inference processes of the MV-Adapter.

4.1 CONDITION GUIDER

We design a general condition guider that supports encoding both camera and geometric representations, enabling T2I models to perform multi-view generation under various guidance.

Camera conditioning. To condition on the camera pose, we use a camera ray representation (“raymap”) that shares the same height and width as the latent representations in the pre-trained T2I models and encodes the ray origin and direction at each spatial location (Watson et al., 2022; Sajjadi et al., 2022; Gao et al., 2024).

Geometry conditioning. Geometry-guided multi-view generation helps applications like texture generation. To condition on the geometry information, we use a global, rather than view-dependent representation that contains position maps and normal maps (Li et al., 2023b; Bensadoun et al., 2024). Each pixel in the position map represents the coordinates of the point on the shape, which provide point correspondences across different views. Normal maps provide orientation information and capture fine geometric details, helping produce detailed textures. We concatenate the position map and normal map along to form a composite geometric conditioning input for each view.

Encoder design. To encode the camera or geometry representation, we design a simple and lightweight condition guider for the conditioning maps c_m ($c_m \in \mathbb{R}^{n \times 6 \times h \times w}$). The condition guider consists of a series of convolutional networks, which contain feature extraction blocks and downsampling layers to adapt the feature resolution to the features in the U-Net encoder. The extracted multi-scale features are then added to the corresponding scales in the U-Net, enabling the

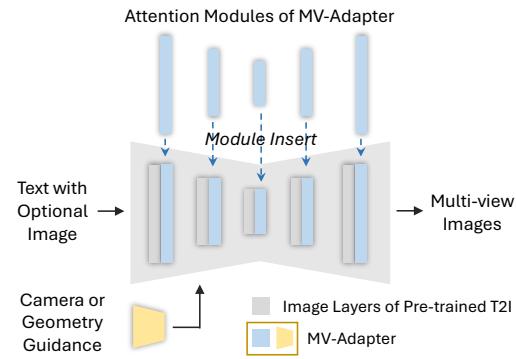


Figure 2: Inference pipeline.

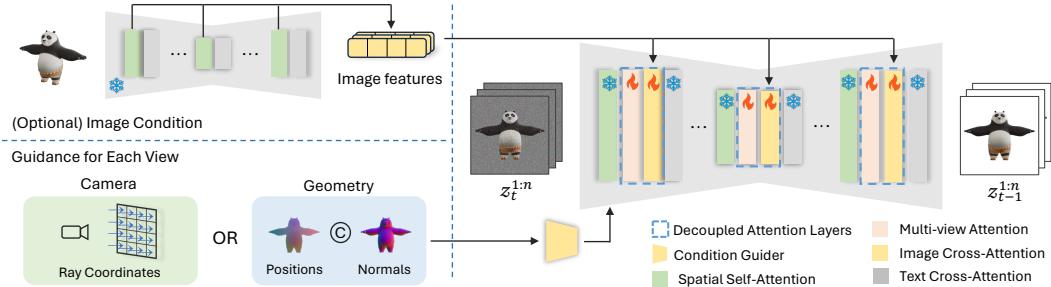


Figure 3: Overview of MV-Adapter. Our MV-Adapter consists of two components: 1) a condition guider that encodes camera or geometry condition; 2) decoupled attention layers that contain multi-view attention for learning multi-view consistency, and optional image cross-attention to support image-conditioned generation, where we use the pre-trained U-Net to encode fine-grained information of the reference image. After training, MV-Adapter can be inserted into any personalized or distilled T2I to generate multi-view images while leveraging the specific strengths of base models.

model to integrate the conditioning information seamlessly at multiple levels. In theory, the input to our encoder is not limited to specific types of conditions; it can also be extended to a wider variety of maps, such as depth maps and pose maps.

4.2 DECOUPLED ATTENTION

We introduce a decoupled attention mechanism, where we retain the original spatial self-attention layers and duplicate them to create new multi-view attention layers as well as image cross-attention layers for image-conditioned generation. These three types of attention layers are organized in a parallel architecture, which ensures that the new attention layers can fully inherit the powerful priors of the pre-trained self-attention layers, thus enabling efficient learning of geometric knowledge.

Duplication of spatial self-attention. Our design adheres to the principle of preserving the original network structure and feature space of the base T2I model. Existing methods like MVDream (Shi et al., 2023b) and Zero123++ (Shi et al., 2023a) modify the base model’s self-attention layers to include multi-view or reference features, which disrupts the learned priors and requires full model fine-tuning. Here we duplicate the structure and weights of spatial self-attention layers to create new multi-view attention and image cross-attention layers, and initialize the output projections of these new attention layers to zero. This allows the new layers to learn geometric knowledge without interfering with the original model, ensuring excellent adaptability.

Parallel attention architecture. In the pre-trained T2I model, the spatial self-attention layer and text cross-attention layer are connected serially through residual connections. Suppose feature f^{in} is the input of the attention block, we can express the process as

$$\begin{aligned} f^{self} &= \text{SelfAttn}(f^{in}) + f^{in}; \\ f^{cross} &= \text{CrossAttn}(f^{self}) + f^{self} \end{aligned} \quad (2)$$

A straightforward method to incorporate new attention layers is to append them after the original layers, connecting them in a serial manner. However, the sequential arrangement may not effectively utilize the image priors modeled by the pre-trained self-attention layers, as it requires the new layers to learn from scratch. Even if we initialize the new layers with the pre-trained weights, the features input to these serially organized layers are in a different domain, causing the initialization to be ineffective. To

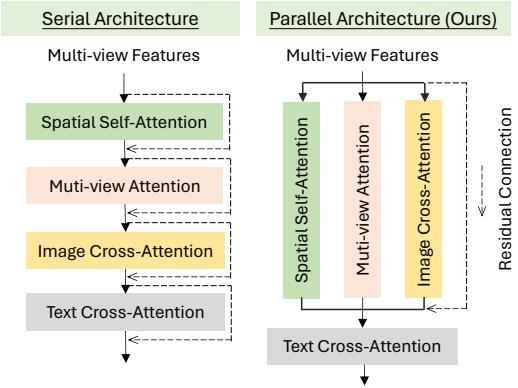


Figure 4: Serial vs parallel architecture.

fully exploit the effective priors of the spatial self-attention layers, we adopt a parallel architecture, as shown in Fig. 4. The process can be formulated as

$$\mathbf{f}^{self} = \text{SelfAttn}(\mathbf{f}^{in}) + \text{MultiViewAttn}(\mathbf{f}^{in}) + \text{ImageCrossAttn}(\mathbf{f}^{in}, \mathbf{f}^{ref}) + \mathbf{f}^{in} \quad (3)$$

where \mathbf{f}^{ref} refers to features of the reference image. Since the features \mathbf{f}^{in} fed into the new layers are the same as those to the self-attention layer, we can effectively initialize them with the pre-trained layers to transfer the image priors. We zero-initialize the output projection layer of the new layers to ensure that the initial output does not disrupt the original feature space. This architectural choice allows the model to build upon the established priors, facilitating efficient learning of multi-view consistency and image-conditioned generation, while preserving the original space of the base T2Is.

Details of multi-view attention. We design different strategies for multi-view attention to meet the specific needs of different applications. For 3D object generation, we enable the model to generate multi-view images at an elevation of 0° and employ row-wise self-attention (Li et al., 2024). For 3D texture generation, considering the view coverage requirements, in addition to the four views evenly at elevation 0° , we add two views from top and bottom. We then perform both row-wise and column-wise self-attention, enabling efficient information exchange among all views. For arbitrary view generation, we employ full self-attention (Shi et al., 2023b) in our multi-view attention layers.

Details of image cross-attention. To condition on reference images \mathbf{c}_i and achieve, we propose a novel method for incorporating detailed information from the image without altering the original feature space of the T2I model. We employ the pre-trained and frozen T2I U-Net as our image encoder. We pass the clear reference image into this frozen U-Net, setting the timestep $t = 0$, and then extract multi-scale features from the spatial self-attention layers. These fine-grained features contain detailed information about the subject and are injected into the denoising U-Net through the decoupled image cross-attention layers. In this way, we leverage the rich representations learned by the pre-trained model, enabling precise control over the generated content.

4.3 TRAINING AND INFERENCE

During training, we only optimize the MV-Adapter, while freezing weights of the pre-trained T2I models. We train MV-Adapter on the dataset with pairs of a reference image, text and n views, using the same training objective as T2I models:

$$\mathcal{L} = \mathbb{E}_{\mathcal{E}(\mathbf{x}_0^{1:n}), \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I}), \mathbf{c}_t, \mathbf{c}_i, \mathbf{c}_m, t} [\|\epsilon - \epsilon_\theta(\mathbf{z}_t^{1:n}, \mathbf{c}_t, \mathbf{c}_i, \mathbf{c}_m, t)\|_2^2] \quad (4)$$

where \mathbf{c}_t , \mathbf{c}_i and \mathbf{c}_m represent texts, reference images and conditioning maps (*i.e.*, camera or geometry conditions) respectively. We randomly zero out the features of the reference image to drop image conditions, enabling classifier-free guidance at inference. Similar to prior work (Blattmann et al., 2023; Hoogeboom et al., 2023), we shift the noise schedule towards high noise levels as we move from the T2Is to the multi-view diffusion model that captures data of higher dimensionality. We shift the log signal-to-noise ratio by $\log(n)$, where n is the number of generated views.

5 EXPERIMENTS

We implemented MV-Adapter on Stable Diffusion V2.1 (SD2.1) and Stable Diffusion XL (SDXL), training a 512×512 adapter for SD2.1 and a 768×768 adapter for SDXL using a subset of the Objaverse dataset (Deitke et al., 2023). Detailed configurations are provided in the Appendix.

5.1 CAMERA-GUIDED MULTI-VIEW GENERATION

Evaluation on community models and extensions. We evaluated MV-Adapter using representative T2Is and extensions, including personalized models (Ruiz et al., 2023; Hu et al., 2021), efficient distilled models (Luo et al., 2023; Lin et al., 2024), and plugins such as ControlNet (Zhang et al., 2023). We present six qualitative results in Fig. 5. More results can be found in the Appendix.

Comparison with baselines. For text-to-multiview generation, we compared our MV-Adapter with MVDream (Shi et al., 2023b) and SPAD (Kant et al., 2024) on 1,000 prompts from the Objaverse dataset. The results are presented in Fig. 6 and Table 1. For image-to-multiview generation,



Figure 5: Results with community models and extensions. Each sample corresponds to a distinct T2I model or extension. Information about the models can be found in the Appendix.

Table 1: Quantitative comparison on camera-guided text-to-multiview generation.

Method	FID↓	IS↑	CLIP Score↑
MVDream	32.15	14.38	31.76
SPAD	48.79	12.04	30.87
Ours (SD2.1)	31.24	15.01	32.04
Ours (SDXL)	29.71	16.38	33.17

Table 2: Quantitative comparison on camera-guided image-to-multiview generation.

Method	PSNR↑	SSIM↑	LPIPS↓
ImageDream	19.280	0.8472	0.1218
Zero123++	20.312	0.8417	0.1205
CRM	20.185	0.8325	0.1247
SV3D	20.042	0.8267	0.1396
Ouroboros3D	20.810	0.8535	0.1193
Era3D	20.890	0.8601	0.1199
Ours (SD2.1)	20.867	0.8695	0.1147
Ours (SDXL)	22.131	0.8816	0.1002

we conduct comparison with ImageDream (Wang & Shi, 2023), Zero123++ (Shi et al., 2023a), CRM (Wang et al., 2024b), SV3D (Voleti et al., 2024), Ouroboros3D (Wen et al., 2024), and Era3D (Li et al., 2024) on the Google Scanned Objects (GSO) dataset (Downs et al., 2022), as results shown in Fig. 7 and Table 2. Experiments indicate that, by preserving the original feature space of T2I models, our MV-Adapter achieves higher visual fidelity and consistency with conditions.

5.2 GEOMETRY-GUIDED MULTI-VIEW GENERATION

Evaluation on community models and extensions. We evaluated our geometry-guided model with T2I derivative models. The results in Fig. 8 demonstrate the adaptability of MV-Adapter in seamlessly integrating with different base models.

Comparison with baselines. We compare our text- and image-conditioned multi-view-based texture generation method (see Sec. 5.4) with four state-of-the-art methods, including TEX-Ture (Richardson et al., 2023), Text2Tex (Chen et al., 2023), Paint3D (Zeng et al., 2024), SyncMVD (Liu et al., 2023b), and FlashTex (Deng et al., 2024). For our image-to-texture model, we used ControlNet (Zhang et al., 2023) to generate reference images conditioned on text and depth maps. As shown in Fig. 10 and Table 3, compared to these project-and-inpaint or synchronized multi-view texturing methods, our approach fine-tunes additional modules to model geometric associations and preserves the generative capabilities of the base T2I model, thereby producing multi-



Figure 6: Qualitative comparison on camera-guided text-to-multiview generation. our MV-Adapter achieves higher visual fidelity and image-text consistency.



Figure 7: Qualitative comparison on camera-guided image-to-multiview generation.

view consistent and high-quality textures. Additionally, testing on a single RTX 4090 GPU revealed that our method achieves faster generation speeds than the others.

5.3 ABLATION STUDY

We conduct ablation studies to evaluate the efficiency and adaptability of our MV-Adapter, as well as the detailed design of the adapter network.

Efficiency. To assess the training efficiency of our adapter design, we conducted comparison with Era3D (Li et al., 2024), which requires full training rather than fine-tuning only adapters like us. As shown in Table 4, when working with SDXL (Podell et al., 2024), our MV-Adapter significantly reduces training costs, facilitating high-resolution multi-view generation based on larger backbones.

Adaptability. We compare MV-Adapter with the full-trained text-to-multiview generation method MVDream (Shi et al., 2023b) regarding compatibility with T2I derivatives. MVDream, which fine-tunes the whole T2I model, cannot be easily replaced with other T2Is; thus, we integrate LoRA (Hu et al., 2021) for our experiments. As shown in Fig. 9, MVDream struggles to generate images that align with the text and style, whereas our MV-Adapter produces high-quality results, demonstrating its superior adaptability.

Parallel attention architecture. To assess the effectiveness of our proposed parallel attention architecture, we conducted ablation studies on image-to-multi-view generation setting. We report the quantitative and qualitative results of using serial or parallel architecture in Table 5 and Fig. 11. The results show that, the serial setting, which cannot leverage the pre-trained image prior, tends to produce artifacts and inconsistent details with the image input. In contrast, our parallel setting produces high-quality and highly consistent results with the reference image.



Figure 9: Qualitative ablation study on the adaptability of MV-Adapter.



Figure 8: Results of geometry-guided text-to-multiview generation with community models.

Table 3: Quantitative comparison on 3D texture generation. FID and KID ($\times 10^{-4}$) are evaluated on multi-view renderings. Our models achieves best texture quality with faster inference.

Method	FID \downarrow	KID \downarrow	Time \downarrow
TEXture	56.44	61.16	90s
Text2Tex	58.43	60.81	421s
Paint3D	44.38	47.06	60s
SyncMVD	36.13	42.28	50s
FlashTex	50.48	56.36	186s
Ours (SD2.1 - Text)	38.19	42.83	18s
Ours (SD2.1 - Image)	33.93	38.73	19s
Ours (SDXL - Text)	32.75	35.18	32s
Ours (SDXL - Image)	27.28	29.47	33s

Table 4: Comparison of training costs with full-tuning methods (batch size set to 1).

Method	Trainable params \downarrow	Memory usage \downarrow	Training speed \uparrow
Era3D (SD2.1)	993M	36G	2.2iter/s
Ours (SD2.1)	127M	17G	3.1iter/s
Era3D (SDXL)	3.1B	>80G	-
Ours (SDXL)	490M	60G	1.05iter/s

Table 5: Quantitative ablation studies on attention architecture.

Method	PSNR \uparrow	SSIM \uparrow	LPIPS \downarrow
Serial (SDXL)	20.687	0.8681	0.1149
Parallel (SDXL)	22.131	0.8816	0.1002

5.4 APPLICATIONS

3D generation. We Follow the existing pipelines (Li et al., 2024) to achieve 3D generation. After generating multi-view images from text or image conditions using MV-Adapter, we use StableNormal (Ye et al., 2024) to generate corresponding normal maps. The multi-view images and normal maps are then fed into NeuS (Wang et al., 2021) to reconstruct the 3D mesh. We conducted comparison on 3D reconstruction with Era3D (Li et al., 2024), which shares a similar pipeline with our method. Results in Table 6 show that our SD2.1-based MV-Adapter is comparable to Era3D, but our SDXL-based model shows significantly higher performance. These findings underline the scalability of MV-Adapter and its ability to leverage the strengths of state-of-the-art T2I models, providing benefits to 3D generation. More results can be found in the Appendix.

Texture generation. We use backprojection and incidence-based weighted blending techniques (Bensadoun et al., 2024) to map the generated multi-view images onto the UV texture map. Despite optimizing view distribution to enhance coverage, some areas may remain uncovered due to occlusions or extreme angles. To address this, we perform view coverage analysis to identify

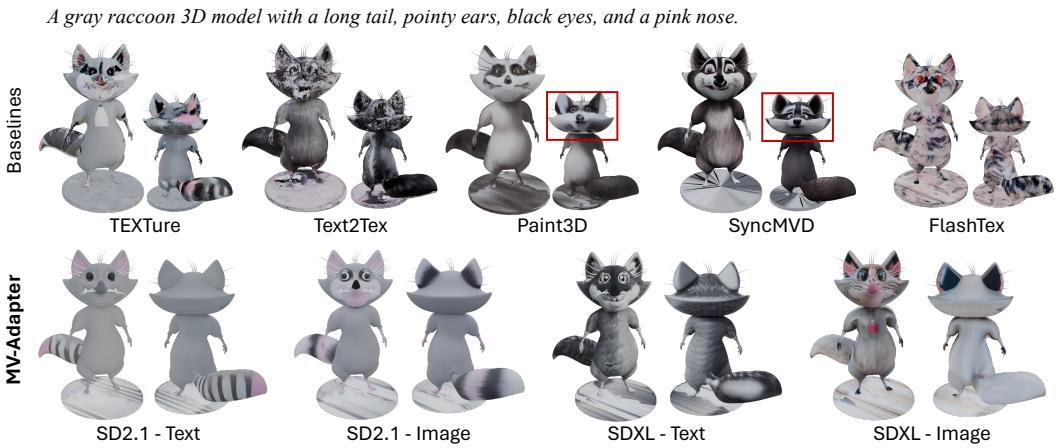


Figure 10: Qualitative comparison on texture generation. We compare our text- and image-conditioned models with baseline methods.

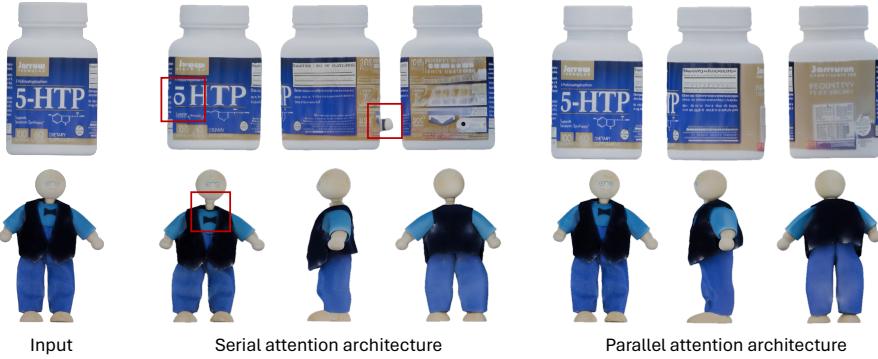


Figure 11: Qualitative ablation study on the attention architecture.

uncovered regions, render images from the current 3D texture for those views, and refine them using an efficient inpainting model (Suvorov et al., 2022). We show more visual results in the Appendix.

Arbitrary view generation. Starting from text or an initial image, we first generate eight anchor views that broadly cover the object. For new target viewpoints, we cluster the viewpoints based on their spatial orientations and select the 4 nearest anchor views to guide the generation of target views. We concatenate these four input views into a single image and input it into the pre-trained T2I U-Net to extract features. Implementation details and visual results are provided in the Appendix and supplementary materials.

6 CONCLUSION

In this paper, we present MV-Adapter, the first adapter-based solution for multi-view image generation. This versatile, plug-and-play adapter enhances text-to-image diffusion models and their derivatives without compromising quality or altering the original feature space. We introduce innovative adapter framework that includes duplicated self-attention layers and a parallel attention architecture, allowing the adapter to efficiently model 3D geometric knowledge. Additionally, we introduced a unified condition encoder that integrates camera parameters and geometric information into spatial map representations, enhancing the model’s versatility and applicability in 3D object generation and texture generation. Extensive evaluations highlight the efficiency, adaptability, and versatility of MV-Adapter across different models and conditions. Overall, MV-Adapter offers an efficient and flexible solution for multi-view image generation, significantly broadening the capabilities of pre-trained T2I models and presenting exciting possibilities for a wide range of applications.

REFERENCES

- Yogesh Balaji, Seungjun Nah, Xun Huang, Arash Vahdat, Jiaming Song, Qinsheng Zhang, Karsten Kreis, Miika Aittala, Timo Aila, Samuli Laine, et al. ediff-i: Text-to-image diffusion models with an ensemble of expert denoisers. *arXiv preprint arXiv:2211.01324*, 2022.
- Raphael Bensadoun, Yanir Kleiman, Idan Azuri, Omri Harosh, Andrea Vedaldi, Natalia Neverova, and Oran Gafni. Meta 3d texturegen: Fast and consistent texture generation for 3d objects. *arXiv preprint arXiv:2407.02430*, 2024.
- Andreas Blattmann, Tim Dockhorn, Sumith Kulal, Daniel Mendelevitch, Maciej Kilian, Dominik Lorenz, Yam Levi, Zion English, Vikram Voleti, Adam Letts, et al. Stable video diffusion: Scaling latent video diffusion models to large datasets. *arXiv preprint arXiv:2311.15127*, 2023.
- Pu Cao, Feng Zhou, Qing Song, and Lu Yang. Controllable generation with text-to-image diffusion models: A survey. *arXiv preprint arXiv:2403.04279*, 2024.
- Dave Zhenyu Chen, Yawar Siddiqui, Hsin-Ying Lee, Sergey Tulyakov, and Matthias Nießner. Text2tex: Text-driven texture synthesis via diffusion models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 18558–18568, 2023.
- Matt Deitke, Dustin Schwenk, Jordi Salvador, Luca Weihs, Oscar Michel, Eli VanderBilt, Ludwig Schmidt, Kiana Ehsani, Aniruddha Kembhavi, and Ali Farhadi. Objaverse: A universe of annotated 3d objects. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13142–13153, 2023.
- Kangle Deng, Timothy Omernick, Alexander Weiss, Deva Ramanan, Jun-Yan Zhu, Tinghui Zhou, and Maneesh Agrawala. Flashtex: Fast relightable mesh texturing with lightcontrolnet. *arXiv preprint arXiv:2402.13251*, 2024.
- Prafulla Dhariwal and Alexander Nichol. Diffusion models beat gans on image synthesis. *Advances in neural information processing systems*, 34:8780–8794, 2021.
- Laura Downs, Anthony Francis, Nate Koenig, Brandon Kinman, Ryan Hickman, Krista Reymann, Thomas B McHugh, and Vincent Vanhoucke. Google scanned objects: A high-quality dataset of 3d scanned household items. In *2022 International Conference on Robotics and Automation (ICRA)*, pp. 2553–2560. IEEE, 2022.
- Rinon Gal, Yuval Alaluf, Yuval Atzmon, Or Patashnik, Amit H Bermano, Gal Chechik, and Daniel Cohen-Or. An image is worth one word: Personalizing text-to-image generation using textual inversion. *arXiv preprint arXiv:2208.01618*, 2022.
- Ruiqi Gao, Aleksander Holynski, Philipp Henzler, Arthur Brussee, Ricardo Martin-Brualla, Pratul Srinivasan, Jonathan T Barron, and Ben Poole. Cat3d: Create anything in 3d with multi-view diffusion models. *arXiv preprint arXiv:2405.10314*, 2024.
- Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. *arXiv preprint arXiv:2207.12598*, 2022.
- Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *Advances in neural information processing systems*, 33:6840–6851, 2020.
- Emiel Hoogeboom, Jonathan Heek, and Tim Salimans. simple diffusion: End-to-end diffusion for high resolution images. In *International Conference on Machine Learning*, pp. 13213–13232. PMLR, 2023.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*, 2021.
- Zehuan Huang, Hao Wen, Junting Dong, Yaohui Wang, Yangguang Li, Xinyuan Chen, Yan-Pei Cao, Ding Liang, Yu Qiao, Bo Dai, et al. Epidiff: Enhancing multi-view synthesis via localized epipolar-constrained diffusion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 9784–9794, 2024.

-
- Yoonwoo Jeong, Jinwoo Lee, Chiheon Kim, Minsu Cho, and Doyup Lee. Nvs-adapter: Plug-and-play novel view synthesis from a single image. In *European Conference on Computer Vision*, pp. 449–466. Springer, 2025.
- Yash Kant, Aliaksandr Siarohin, Ziyi Wu, Michael Vasilkovsky, Guocheng Qian, Jian Ren, Riza Alp Guler, Bernard Ghanem, Sergey Tulyakov, and Igor Gilitschenski. Spad: Spatially aware multi-view diffusers. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 10026–10038, 2024.
- Nupur Kumari, Bingliang Zhang, Richard Zhang, Eli Shechtman, and Jun-Yan Zhu. Multi-concept customization of text-to-image diffusion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 1931–1941, 2023.
- Black Forest Labs. Flux. [Online], 2024. <https://github.com/black-forest-labs/flux>.
- Jiahao Li, Hao Tan, Kai Zhang, Zexiang Xu, Fujun Luan, Yinghao Xu, Yicong Hong, Kalyan Sunkavalli, Greg Shakhnarovich, and Sai Bi. Instant3d: Fast text-to-3d with sparse-view generation and large reconstruction model. *arXiv preprint arXiv:2311.06214*, 2023a.
- Peng Li, Yuan Liu, Xiaoxiao Long, Feihu Zhang, Cheng Lin, Mengfei Li, Xingqun Qi, Shanghang Zhang, Wenhan Luo, Ping Tan, et al. Era3d: High-resolution multiview diffusion using efficient row-wise attention. *arXiv preprint arXiv:2405.11616*, 2024.
- Weiyu Li, Rui Chen, Xuelin Chen, and Ping Tan. Sweetdreamer: Aligning geometric priors in 2d diffusion for consistent text-to-3d. *arXiv preprint arXiv:2310.02596*, 2023b.
- Shanchuan Lin, Anran Wang, and Xiao Yang. Sdxl-lightning: Progressive adversarial diffusion distillation. *arXiv preprint arXiv:2402.13929*, 2024.
- Yuan Liu, Cheng Lin, Zijiao Zeng, Xiaoxiao Long, Lingjie Liu, Taku Komura, and Wenping Wang. Syncdreamer: Generating multiview-consistent images from a single-view image. *arXiv preprint arXiv:2309.03453*, 2023a.
- Yuxin Liu, Minshan Xie, Hanyuan Liu, and Tien-Tsin Wong. Text-guided texturing by synchronized multi-view diffusion. *arXiv preprint arXiv:2311.12891*, 2023b.
- Xiaoxiao Long, Yuan-Chen Guo, Cheng Lin, Yuan Liu, Zhiyang Dou, Lingjie Liu, Yuexin Ma, Song-Hai Zhang, Marc Habermann, Christian Theobalt, et al. Wonder3d: Single image to 3d using cross-domain diffusion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 9970–9980, 2024.
- Simian Luo, Yiqin Tan, Longbo Huang, Jian Li, and Hang Zhao. Latent consistency models: Synthesizing high-resolution images with few-step inference. *arXiv preprint arXiv:2310.04378*, 2023.
- Tiange Luo, Chris Rockwell, Honglak Lee, and Justin Johnson. Scalable 3d captioning with pre-trained models. *Advances in Neural Information Processing Systems*, 36, 2024.
- Jian Ma, Junhao Liang, Chen Chen, and Haonan Lu. Subject-diffusion: Open domain personalized text-to-image generation without test-time fine-tuning. In *ACM SIGGRAPH 2024 Conference Papers*, pp. 1–12, 2024.
- Chenlin Meng, Robin Rombach, Ruiqi Gao, Diederik Kingma, Stefano Ermon, Jonathan Ho, and Tim Salimans. On distillation of guided diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 14297–14306, 2023.
- Ron Mokady, Amir Hertz, Kfir Aberman, Yael Pritch, and Daniel Cohen-Or. Null-text inversion for editing real images using guided diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 6038–6047, 2023.
- Chong Mou, Xintao Wang, Liangbin Xie, Yanze Wu, Jian Zhang, Zhongang Qi, and Ying Shan. T2i-adapter: Learning adapters to dig out more controllable ability for text-to-image diffusion models. *Proceedings of the AAAI Conference on Artificial Intelligence*, 38(5):4296–4304, Mar. 2024. doi: 10.1609/aaai.v38i5.28226. URL <https://ojs.aaai.org/index.php/AAAI/article/view/28226>.

-
- Alexander Quinn Nichol, Prafulla Dhariwal, Aditya Ramesh, Pranav Shyam, Pamela Mishkin, Bob McGrew, Ilya Sutskever, and Mark Chen. GLIDE: towards photorealistic image generation and editing with text-guided diffusion models. In *International Conference on Machine Learning, ICML 2022, 17-23 July 2022, Baltimore, Maryland, USA*, volume 162 of *Proceedings of Machine Learning Research*, pp. 16784–16804, 2022.
- William Peebles and Saining Xie. Scalable diffusion models with transformers. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 4195–4205, 2023.
- Dustin Podell, Zion English, Kyle Lacey, Andreas Blattmann, Tim Dockhorn, Jonas Müller, Joe Penna, and Robin Rombach. SDXL: improving latent diffusion models for high-resolution image synthesis. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*, 2024.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *International conference on machine learning*, pp. 8748–8763. PMLR, 2021.
- Aditya Ramesh, Mikhail Pavlov, Gabriel Goh, Scott Gray, Chelsea Voss, Alec Radford, Mark Chen, and Ilya Sutskever. Zero-shot text-to-image generation. In *International conference on machine learning*, pp. 8821–8831. Pmlr, 2021.
- Aditya Ramesh, Prafulla Dhariwal, Alex Nichol, Casey Chu, and Mark Chen. Hierarchical text-conditional image generation with clip latents. *arXiv preprint arXiv:2204.06125*, 1(2):3, 2022.
- Elad Richardson, Gal Metzer, Yuval Alaluf, Raja Giryes, and Daniel Cohen-Or. Texture: Text-guided texturing of 3d shapes. In *ACM SIGGRAPH 2023 conference proceedings*, pp. 1–11, 2023.
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 10684–10695, 2022.
- Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *Medical image computing and computer-assisted intervention-MICCAI 2015: 18th international conference, Munich, Germany, October 5-9, 2015, proceedings, part III 18*, pp. 234–241. Springer, 2015.
- Nataniel Ruiz, Yuanzhen Li, Varun Jampani, Yael Pritch, Michael Rubinstein, and Kfir Aberman. Dreambooth: Fine tuning text-to-image diffusion models for subject-driven generation. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 22500–22510, 2023.
- Chitwan Saharia, William Chan, Saurabh Saxena, Lala Li, Jay Whang, Emily L Denton, Kamyar Ghasemipour, Raphael Gontijo Lopes, Burcu Karagol Ayan, Tim Salimans, et al. Photorealistic text-to-image diffusion models with deep language understanding. *Advances in neural information processing systems*, 35:36479–36494, 2022.
- Mehdi SM Sajjadi, Henning Meyer, Etienne Pot, Urs Bergmann, Klaus Greff, Noha Radwan, Suhani Vora, Mario Lučić, Daniel Duckworth, Alexey Dosovitskiy, et al. Scene representation transformer: Geometry-free novel view synthesis through set-latent scene representations. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 6229–6238, 2022.
- Viraj Shah, Nataniel Ruiz, Forrester Cole, Erika Lu, Svetlana Lazebnik, Yuanzhen Li, and Varun Jampani. Ziplora: Any subject in any style by effectively merging loras. *arXiv preprint arXiv:2311.13600*, 2023.
- Jing Shi, Wei Xiong, Zhe Lin, and Hyun Joon Jung. Instantbooth: Personalized text-to-image generation without test-time finetuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 8543–8552, 2024.

-
- Ruoxi Shi, Hansheng Chen, Zhuoyang Zhang, Minghua Liu, Chao Xu, Xinyue Wei, Linghao Chen, Chong Zeng, and Hao Su. Zero123++: a single image to consistent multi-view diffusion base model. *arXiv preprint arXiv:2310.15110*, 2023a.
- Yichun Shi, Peng Wang, Jianglong Ye, Mai Long, Kejie Li, and Xiao Yang. Mvdream: Multi-view diffusion for 3d generation. *arXiv preprint arXiv:2308.16512*, 2023b.
- Yukai Shi, Jianan Wang, He Cao, Boshi Tang, Xianbiao Qi, Tianyu Yang, Yukun Huang, Shilong Liu, Lei Zhang, and Heung-Yeung Shum. Toss: High-quality text-guided novel view synthesis from a single image. *arXiv preprint arXiv:2310.10644*, 2023c.
- Jiaming Song, Chenlin Meng, and Stefano Ermon. Denoising diffusion implicit models. *arXiv preprint arXiv:2010.02502*, 2020.
- Kunpeng Song, Yizhe Zhu, Bingchen Liu, Qing Yan, Ahmed Elgammal, and Xiao Yang. Moma: Multimodal lilm adapter for fast personalized image generation. *arXiv preprint arXiv:2404.05674*, 2024.
- Yang Song, Prafulla Dhariwal, Mark Chen, and Ilya Sutskever. Consistency models. *arXiv preprint arXiv:2303.01469*, 2023.
- Roman Suvorov, Elizaveta Logacheva, Anton Mashikhin, Anastasia Remizova, Arsenii Ashukha, Aleksei Silvestrov, Naejin Kong, Harshith Goka, Kiwoong Park, and Victor Lempitsky. Resolution-robust large mask inpainting with fourier convolutions. In *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, pp. 2149–2159, 2022.
- Shitao Tang, Fuyang Zhang, Jiacheng Chen, Peng Wang, and Yasutaka Furukawa. Mvdiffusion: Enabling holistic multi-view image generation with correspondence-aware diffusion. *arXiv*, 2023.
- Shitao Tang, Jiacheng Chen, Dilin Wang, Chengzhou Tang, Fuyang Zhang, Yuchen Fan, Vikas Chandra, Yasutaka Furukawa, and Rakesh Ranjan. Mvdiffusion++: A dense high-resolution multi-view diffusion model for single or sparse-view 3d object reconstruction. *arXiv preprint arXiv:2402.12712*, 2024.
- Kolors Team. Kolors: Effective training of diffusion model for photorealistic text-to-image synthesis. *arXiv preprint*, 2024.
- Vikram Voleti, Chun-Han Yao, Mark Boss, Adam Letts, David Pankratz, Dmitry Tochilkin, Christian Laforte, Robin Rombach, and Varun Jampani. Sv3d: Novel multi-view synthesis and 3d generation from a single image using latent video diffusion. *arXiv preprint arXiv:2403.12008*, 2024.
- Haofan Wang, Qixun Wang, Xu Bai, Zekui Qin, and Anthony Chen. Instantstyle: Free lunch towards style-preserving in text-to-image generation. *arXiv preprint arXiv:2404.02733*, 2024a.
- Peng Wang and Yichun Shi. Imagedream: Image-prompt multi-view diffusion for 3d generation. *arXiv preprint arXiv:2312.02201*, 2023.
- Peng Wang, Lingjie Liu, Yuan Liu, Christian Theobalt, Taku Komura, and Wenping Wang. Neus: Learning neural implicit surfaces by volume rendering for multi-view reconstruction. *arXiv preprint arXiv:2106.10689*, 2021.
- Zhengyi Wang, Yikai Wang, Yifei Chen, Chendong Xiang, Shuo Chen, Dajiang Yu, Chongxuan Li, Hang Su, and Jun Zhu. Crm: Single image to 3d textured mesh with convolutional reconstruction model. *arXiv preprint arXiv:2403.05034*, 2024b.
- Daniel Watson, William Chan, Ricardo Martin-Brualla, Jonathan Ho, Andrea Tagliasacchi, and Mohammad Norouzi. Novel view synthesis with diffusion models. *arXiv preprint arXiv:2210.04628*, 2022.
- Hao Wen, Zehuan Huang, Yaohui Wang, Xinyuan Chen, Yu Qiao, and Lu Sheng. Ouroboros3d: Image-to-3d generation via 3d-aware recursive diffusion. *arXiv preprint arXiv:2406.03184*, 2024.

-
- Zhuoyi Yang, Jiayan Teng, Wendi Zheng, Ming Ding, Shiyu Huang, Jiazheng Xu, Yuanming Yang, Wenyi Hong, Xiaohan Zhang, Guanyu Feng, et al. Cogvideox: Text-to-video diffusion models with an expert transformer. *arXiv preprint arXiv:2408.06072*, 2024.
- Chongjie Ye, Lingteng Qiu, Xiaodong Gu, Qi Zuo, Yushuang Wu, Zilong Dong, Liefeng Bo, Yuliang Xiu, and Xiaoguang Han. Stablenormal: Reducing diffusion variance for stable and sharp normal. *ACM Transactions on Graphics (TOG)*, 2024.
- Hu Ye, Jun Zhang, Sibo Liu, Xiao Han, and Wei Yang. Ip-adapter: Text compatible image prompt adapter for text-to-image diffusion models. *arXiv preprint arXiv:2308.06721*, 2023.
- Xianggang Yu, Mutian Xu, Yidan Zhang, Haolin Liu, Chongjie Ye, Yushuang Wu, Zizheng Yan, Chenming Zhu, Zhangyang Xiong, Tianyou Liang, et al. Mvimgnet: A large-scale dataset of multi-view images. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 9150–9161, 2023.
- Xianfang Zeng, Xin Chen, Zhongqi Qi, Wen Liu, Zibo Zhao, Zhibin Wang, Bin Fu, Yong Liu, and Gang Yu. Paint3d: Paint anything 3d with lighting-less texture diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4252–4262, 2024.
- Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image diffusion models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 3836–3847, 2023.
- Chuanxia Zheng and Andrea Vedaldi. Free3d: Consistent novel view synthesis without 3d representation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 9720–9731, 2024.
- Zangwei Zheng, Xiangyu Peng, Tianji Yang, Chenhui Shen, Shenggui Li, Hongxin Liu, Yukun Zhou, Tianyi Li, and Yang You. Open-sora: Democratizing efficient video production for all, March 2024. URL <https://github.com/hpcatech/Open-Sora>.

A APPENDIX

A.1 BACKGROUND

Stable Diffusion (SD) and Stable Diffusion XL (SDXL). We adopt Stable Diffusion (Rombach et al., 2022) and Stable Diffusion XL (Podell et al., 2024) as our base T2I models, since they have a well-developed community with many powerful derivatives for evaluation. SD and SDXL perform the diffusion process within the latent space of a pre-trained autoencoder $\mathcal{E}(\cdot)$ and $\mathcal{D}(\cdot)$. In training, an encoded image $z_0 = \mathcal{E}(x_0)$ is perturbed to z_t at step t by the forward diffusion. The denoising network ϵ_θ learns to reverse this process by predicting the added noise, encouraged by an MSE loss:

$$\mathcal{L} = \mathbb{E}_{\mathcal{E}(x_0), \epsilon \sim \mathcal{N}(0, I), c, t} [\|\epsilon - \epsilon_\theta(z_t, c, t)\|_2^2] \quad (5)$$

where c denotes the conditioning texts. In SD, ϵ_θ is implemented as a UNet (Ronneberger et al., 2015) consisting of pairs of down/up sample blocks and a middle block. Each block contains pairs of spatial self-attention layers and cross-attention layers, which are serially connected using the residual structure. SDXL leverages a three times larger UNet backbone than SD for high-resolution image synthesis, and introduces a refinement denoiser to improve the visual fidelity.

A.2 IMPLEMENTATION DETAILS

Dataset. We trained MV-Adapter on a filtered high-quality subset of the Objaverse dataset (Deitke et al., 2023), comprising approximately 70,000 samples, with captions from Cap3D (Luo et al., 2024). To accommodate the efficient multi-view self-attention mechanism, we rendered orthographic views to train the model to generate $n = 6$ views per sample. For the camera-guided generation, we rendered views of 3D models with the elevation angle set to 0° and azimuth angles at $\{0^\circ, 45^\circ, 90^\circ, 180^\circ, 270^\circ, 315^\circ\}$. This distribution aligns with the setting used in Era3D (Li et al., 2024), facilitating the application of a similar image-to-3D pipeline for 3D generation tasks. For the geometry-guided generation, we included four views at an elevation of 0° with azimuth angles

of $\{0^\circ, 90^\circ, 180^\circ, 270^\circ\}$, added two additional views from the top and bottom. In addition to the target views, we rendered five random views within a certain frontal range of the models to serve as reference images during training.

Training. We utilized two versions of Stable Diffusion (Rombach et al., 2022) as the base models for training. Specifically, we trained a 512-resolution model based on Stable Diffusion 2.1 (SD2.1) and a 768-resolution model based on Stable Diffusion XL (SDXL). During training, we randomly dropped the text condition with a probability of 0.1, the image condition with a probability of 0.1, and both text and image conditions simultaneously with a probability of 0.1. Following prior work (Hoogeboom et al., 2023; Blattmann et al., 2023), we shifted the noise schedule to higher noise levels by adjusting the log signal-to-noise ratio (SNR) by $\log(n)$, where $n = 6$ is the number of the generated views. For the specific training configurations, we used a learning rate of 5×10^{-5} and trained the MV-Adapter on 8 NVIDIA A100 GPUs for 10 epochs.

Inference. In our experimental setup, we generated multi-view images using the DDPM sampler (Ho et al., 2020) with classifier-free guidance (Ho & Salimans, 2022), and set the number of inference steps to 50. For generation conditioned solely on text (*i.e.*, setting the weight of the image condition λ_i to 0), we set the guidance scale to 7.0. For image-conditioned generation, we set the guidance scale of image condition α and text condition β to 3.0. Following TOSS (Shi et al., 2023c), the calculation can be expressed as:

$$\begin{aligned} \hat{\epsilon}_\theta(\mathbf{z}_t^{1:n}, \mathbf{c}_t, \mathbf{c}_i, \mathbf{c}_m, t) &= \epsilon_\theta(\mathbf{z}_t^{1:n}, \emptyset, \emptyset, \mathbf{c}_m, t) \\ &\quad + \alpha [\epsilon_\theta(\mathbf{z}_t^{1:n}, \emptyset, \mathbf{c}_i, \mathbf{c}_m, t) - \epsilon_\theta(\mathbf{z}_t^{1:n}, \emptyset, \emptyset, \mathbf{c}_m, t)] \\ &\quad + \beta [\epsilon_\theta(\mathbf{z}_t^{1:n}, \mathbf{c}_t, \mathbf{c}_i, \mathbf{c}_m, t) - \epsilon_\theta(\mathbf{z}_t^{1:n}, \emptyset, \mathbf{c}_i, \mathbf{c}_m, t)] \end{aligned} \quad (6)$$

where \mathbf{c}_t , \mathbf{c}_i and \mathbf{c}_m represent texts, reference images and conditioning maps (*i.e.*, camera or geometry conditions) respectively. Since we did not drop \mathbf{c}_m during the training process, we do not use the classifier-free guidance method for it.

Comparison with baselines. We conducted comprehensive comparisons with baseline methods across three settings: text-to-multiview generation, image-to-multiview generation, and texture generation. In these experiments, we evaluated both versions of MV-Adapter based on Stable Diffusion 2.1 (SD2.1) (Rombach et al., 2022) and Stable Diffusion XL (SDXL) (Podell et al., 2024), demonstrating the performance gains brought by MV-Adapter due to its efficient training and scalability.

For text-to-multiview generation, we selected MVDream (Shi et al., 2023b) and SPAD (Kant et al., 2024) as baseline methods. MVDream extends the original self-attention mechanism of T2I models to the multi-view domain. SPAD introduces epipolar constraints into the multi-view attention mechanism. We tested on 1,000 prompts selected from the Objaverse dataset (Deitke et al., 2023). We computed Fréchet Inception Distance (FID), Inception Score (IS), and CLIP Score on all generated views to assess the quality of the generated images and their alignment with the textual prompts.

For image-to-multiview generation, we compared our method with ImageDream (Wang & Shi, 2023), Zero123++(Shi et al., 2023a), CRM(Wang et al., 2024b), SV3D (Voleti et al., 2024), Ouroboros3D (Wen et al., 2024), and Era3D (Li et al., 2024). ImageDream, Zero123++, CRM, and Era3D generally fall into the category of modifying the original network architecture of T2I models to extend them for multi-view generation. SV3D and Ouroboros3D fine-tune text-to-video (T2V) models to achieve multi-view generation. We selected 100 assets covering multiple object categories from the Google Scanned Objects (GSO) dataset (Downs et al., 2022) as our test set. For each asset, we rendered input images from front-facing views, with input views randomly distributed in azimuth angles between -45° and 45° and elevation angles between -10° and 30° . We evaluated the generated multi-view images by computing Peak Signal-to-Noise Ratio (PSNR), Structural Similarity Index Measure (SSIM), and Learned Perceptual Image Patch Similarity (LPIPS) between the generated images and the ground truth, assessing both the consistency and quality of the outputs.

For 3D texture generation, we compared our text-based and image-based models with project-and-paint methods such as TEXTure (Richardson et al., 2023), Text2Tex (Chen et al., 2023), and Paint3D (Zeng et al., 2024), the synchronized multi-view texturing method SyncMVD (Liu et al., 2023b), and the optimization-based method FlashTex (Deng et al., 2024). We randomly selected 200 models along with their captions from the Objaverse (Deitke et al., 2023) dataset for testing.

Table 7: Community models and extensions for evaluation.

Category	Model Name	Domain	Model Type
Personalized T2I	Dreamshaper ¹	General	T2I Base Model
	RealVisXL ²	Realistic	T2I Base Model
	Animagine-xl ³	2D Cartoon	T2I Base Model
	3D Render Style XL ⁴	3D Cartoon	LoRA
	Pokemon Trainer Sprite PixelArt ⁵	Pixel Art	LoRA
	Chalk Sketch SDXL ⁶	Chalk Sketch	LoRA
	Chinese Ink LoRA ⁷	Color Ink	LoRA
	Zen Ink Wash Sumi-e ⁸	Wash Ink	LoRA
	Watercolor Style SDXL ⁹	Watercolor	LoRA
	Papercut SDXL ¹⁰	Papercut	LoRA
	Furry Enhancer ¹¹	Enhancer	LoRA
	White Pitbull Dog SDXL ¹²	Concept	LoRA
	Spider spirit fourth sister ¹³	Concept	LoRA
Distilled T2I	SDXL-Lightning ¹⁴	Few Step	T2I Base Model
	LCM-SDXL ¹⁵	Few Step	T2I Base Model
Extension	ControlNet Openpose ¹⁶	Spatial Control	Plugin
	ControlNet Scribble ¹⁷	Spatial Control	Plugin
	ControlNet Tile ¹⁸	Image Deblur	Plugin
	T2I-Adapter Sketch ¹⁹	Spatial Control	Plugin
	IP-Adapter ²⁰	Image Prompt	Plugin

Multiple views were rendered from the generated 3D textures, and we computed FID and Kernel Inception Distance (KID) of them to evaluate the quality of the generated textures. Additionally, we recorded the texture generation time to assess the inference efficiency of each method.

Community models and extensions for evaluation. To ensure a comprehensive benchmark, we selected a diverse set of representative T2I derivative models and extensions from the community for evaluation. As illustrated in Table 7, these models include personalized models that encompass various domains such as anime, stylistic paintings, and realistic photographic images, as well as efficient distilled models and plugins for controllable generation. They cover a wide range of subjects,

¹<https://civitai.com/models/112902?modelVersionId=126688>

²<https://civitai.com/models/139562?modelVersionId=789646>

³<https://huggingface.co/cagliostrolab/animagine-xl-3.1>

⁴https://huggingface.co/goofyai/3d_render_style_xl

⁵<https://civitai.com/models/159333/pokemon-trainer-sprite-pixelart?modelVersionId=443092>

⁶<https://huggingface.co/JerryOrbachJr/Chalk-Sketch-SDXL>

⁷https://huggingface.co/ming-yang/sdxl_chinese_ink_lora

⁸<https://civitai.com/models/647926/zen-ink-wash-sumi-e-sdxl-pony-flux?modelVersionId=724876>

⁹<https://civitai.com/models/484723/watercolor-style-sdxl>

¹⁰https://huggingface.co/TheLastBen/Papercut_SDXL

¹¹<https://civitai.com/models/310964/furry-enhancer?modelVersionId=558568>

¹²<https://civitai.com/models/700883/white-pitbull-dog-sdxl?modelVersionId=787948>

¹³<https://civitai.com/models/689010/pony-black-myth-wukong-spider-spirit-fourth-sister?modelVersionId=771146>

¹⁴<https://huggingface.co/ByteDance/SDXL-Lightning>

¹⁵<https://huggingface.co/latent-consistency/lcm-sdxl>

¹⁶<https://huggingface.co/xinsir/controlnet-openpose-sdxl-1.0>

¹⁷<https://huggingface.co/xinsir/controlnet-scribble-sdxl-1.0>

¹⁸<https://huggingface.co/xinsir/controlnet-tile-sdxl-1.0>

¹⁹<https://huggingface.co/TencentARC/t2i-adapter-sketch-sdxl-1.0>

²⁰<https://huggingface.co/h94/IP-Adapter>

including portraits, animals, landscapes, and more. This selection enables a thorough evaluation of our approach across different styles and content, demonstrating the adaptability and generality of MV-Adapter in working with various T2I derivatives and extensions.

A.3 ADDITIONAL DISCUSSIONS

A.3.1 MV-ADAPTER VS. MULTI-VIEW LORA

LoRA (Low-Rank Adaptation) (Hu et al., 2021) offers an alternative approach to achieving plug-and-play multi-view generation. Specifically, using a condition encoder to inject camera representations, we extend the original self-attention mechanism to operate across all pixels of multiple views. During training, we introduce trainable LoRA layers into the network, allowing these layers to learn multi-view consistency or, optionally, generate images conditioned on a reference view. This approach requires the spatial self-attention mechanism to simultaneously capture spatial image knowledge, ensure multi-view consistency, and align generated images with reference views.

However, the multi-view LoRA approach has a notable limitation. The “incremental changes” it introduces to the network are **not orthogonal or decoupled** from those induced by T2I derivatives, such as personalized T2I models or LoRAs. Specifically, layers fine-tuned by multi-view LoRA and those tuned by personalized LoRA often overlap. Note that each weight matrix learned by both represents a linear transformation defined by its columns, so it is intuitive that the merger would retain the information available in these columns only when the columns that are being added are orthogonal to each other (Shah et al., 2023). Clearly, the multi-view LoRA and personalized models are not orthogonal, which often leads to challenges in retaining both sets of learned knowledge. This can result in a trade-off where either multi-view consistency or the fidelity of concepts (such as style or subject identity) is compromised.

In contrast, our proposed **decoupled** attention mechanism encourages different attention layers to specialize in their respective tasks without needing to fine-tune the original spatial self-attention layers. In this design, the layers we train do not overlap with those in the original T2I model, thereby better preserving the original feature space and enhancing compatibility with other models.

We conducted a series of experiments to test these approaches. We trained two versions of multi-view LoRA, targeting different modules: (1) inserting LoRA layers only into the attention layers, and (2) inserting LoRA layers into multiple layers, including the convolutional layers, down-sampling, up-sampling layers, etc. For both settings, we set the LoRA rank to 64 and alpha to 32. As shown in Fig. 12 and Fig. 13, while the multi-view LoRA approach can generate multi-view consistent images when the base model is not changed, it often struggles to maintain multi-view consistency when switching to a different base model or when integrating a new LoRA. In contrast, as demonstrated in Fig. 14, our MV-Adapter, equipped with the decoupled attention mechanism, maintains consistent multi-view generation even when used with personalized models.

Compared to the LoRA mechanism, our decoupled attention-based approach proves more robust and adaptable for extending T2I models to multi-view generation, offering greater flexibility and compatibility with various pre-trained models.

A.3.2 IMAGE RESTORATION CAPABILITIES

During the training of MV-Adapter, we probabilistically compress the resolution of reference images in the training data pairs to enhance the robustness of multi-view generation from images. We observed that the model trained with this approach is capable of generating high-resolution, detailed multi-view images even when the input is low-resolution, as depicted in Fig. 15. Through such training strategy, MV-Adapter has inherent image restoration capabilities and automatically enhances and refines input images during the generation process.

A.3.3 APPLICABILITY OF MV-ADAPTER

Beyond the demonstrated applications in 3D object generation and 3D texture mapping, the MV-Adapter’s strong adaptability and versatility open up a wide array of potential uses in image creation and personalization. For instance, creators can integrate MV-Adapter with their personalized T2I models—customized for specific identities or artistic styles—to generate multi-view images

(Base model: SDXL) Daenerys Targaryen from game of throne, full body, blender 3d, art station



(Base model: Animage-xl) 1 girl, pink hair, pink shirts, smile, shy, masterpiece



(LoRA: Watercolor Style) painting, Burmese Cat, wearing ral-wtrclr, Comic book art



Figure 12: Results of multi-view LoRA (set target modules to attention layers). The azimuth angles of the images from left to right are $0^\circ, 45^\circ, 90^\circ, 180^\circ, 270^\circ, 315^\circ$, corresponding to the front, front-left, left, back, right, and front-right of the object.

(Base model: SDXL) Daenerys Targaryen from game of throne, full body, blender 3d, art station



(Base model: Animage-xl) 1 girl, pink hair, pink shirts, smile, shy, masterpiece



(LoRA: Watercolor Style) painting, Burmese Cat, wearing ral-wtrclr, Comic book art



Figure 13: Results of multi-view LoRA (set target modules to attention layers, convolutional layers, etc.). The azimuth angles of the images from left to right are $0^\circ, 45^\circ, 90^\circ, 180^\circ, 270^\circ, 315^\circ$, corresponding to the front, front-left, left, back, right, and front-right of the object.

that capture consistent perspectives of their unique concepts. Additionally, MV-Adapter can facilitate tasks like multi-view portrait generation, where a subject’s face is rendered consistently across different angles, or stylized multi-view illustrations that maintain artistic coherence across diverse perspectives.

(Base model: SDXL) Daenerys Targaryen from game of thrones, full body, blender 3d, art station



(Base model: Animage-xl) 1 girl, pink hair, pink shirts, smile, shy, masterpiece



(LoRA: Watercolor Style) painting, Burmese Cat, wearing ral-wtrclr; Comic book art

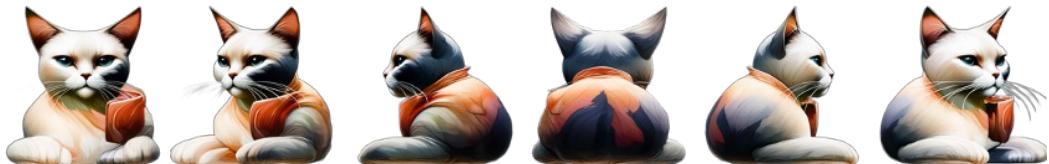


Figure 14: Results of MV-Adapter, which introduces decoupled attention mechanism rather than LoRA. The azimuth angles of the images from left to right are $0^\circ, 45^\circ, 90^\circ, 180^\circ, 270^\circ, 315^\circ$, corresponding to the front, front-left, left, back, right, and front-right of the object.

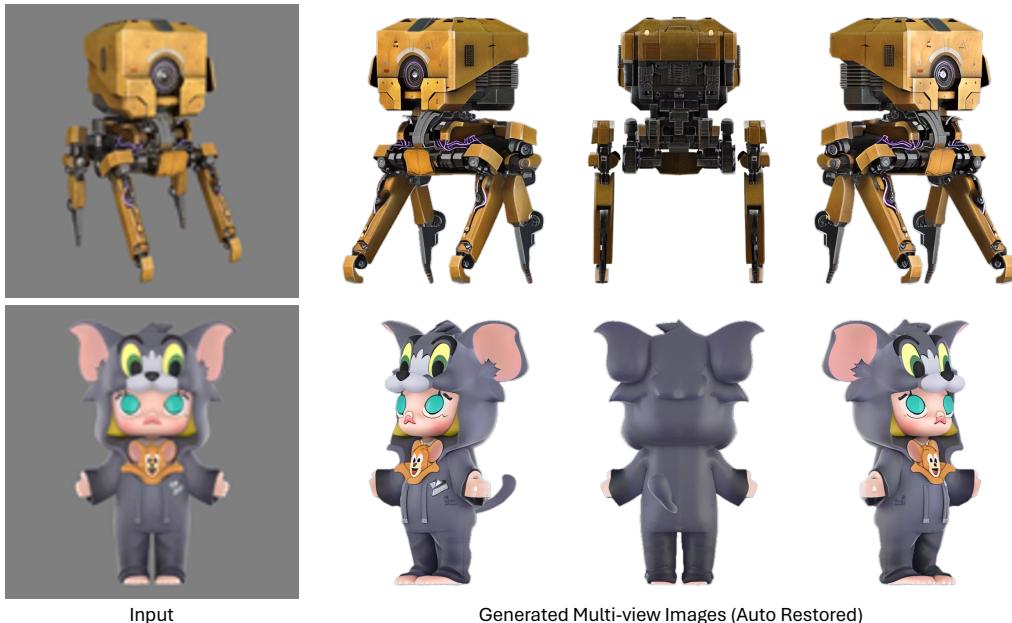


Figure 15: Results on camera-guided image-to-multiview generation with low-resolution images as input.

A.3.4 EXTENDING MV-ADAPTER FOR ARBITRARY VIEW SYNTHESIS

In the main text, we introduced a novel adapter architecture—comprising parallel attention layers and a unified condition encoder—to achieve multi-view generation. We implemented efficient row-wise and column-wise attention mechanisms tailored for two specific applications: 3D object gener-

ation and 3D texture mapping, generating six views accordingly. However, our adapter framework is not limited to these configurations and can be extended to perform arbitrary view synthesis. To explore this capability, we designed a corresponding approach and conducted experiments, training a new version of MV-Adapter to handle arbitrary viewpoints.

Following CAT3D (Gao et al., 2024), we perform multiple rounds of multi-view generation, with the number of views generated each time set to $n = 8$. Starting from text or an initial single image as input, we first generate eight anchor views that broadly cover the object. In practice, these anchor views are positioned at elevations of 0° and 30° , with azimuth angles evenly distributed around the circle (*e.g.* every 45°). For generating new target views, we cluster the viewpoints based on their spatial orientations, grouping them into clusters of 8. We then select the 4 nearest known views from the already generated anchor views to serve as conditions guiding the generation of each target view.

In terms of implementation, the overall framework of our MV-Adapter remains unchanged. We adjust its inputs and specific attention components to accommodate arbitrary view synthesis. First, we set the number of input images to either 1 or 4. When using four input views, we concatenate them into a long image and input this into the pre-trained T2I U-Net to extract features. This simple yet effective method allows the images from the four views to interact within the pre-trained U-Net without requiring additional camera embeddings to represent these views. Second, we utilize full self-attention in the multi-view attention component, expanding the attention scope to enable the generation of target views with more flexible distributions.

To train an MV-Adapter capable of generating arbitrary viewpoints, we rendered data from 40 different views, with elevations of $-10^\circ, 0^\circ, 10^\circ, 20^\circ, 30^\circ$, and azimuth angles evenly distributed around 360 degrees at each elevation layer. We trained the model for 16 epochs. During the first 8 epochs, the model was trained using a setting of one conditional view and eight target anchor views. In the subsequent 8 epochs, we trained with an equal mixture of one condition plus eight target views and four conditions plus eight target views.

As shown in Fig. 16, the visualization results demonstrate that MV-Adapter can generate consistent, high-quality multi-view images beyond the six views designed for specific applications. This extension further verifies the scalability and practicality of our adapter framework, showcasing its potential for arbitrary view synthesis in diverse applications. More results can be found in the supplementary materials.

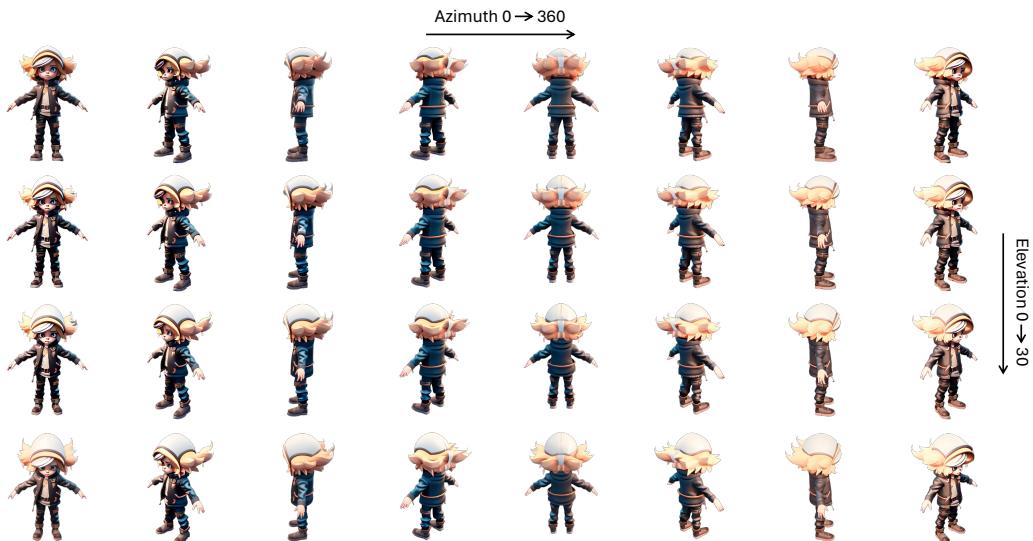


Figure 16: Visualization results using MV-Adapter to generate arbitrary viewpoints.

A.4 LIMITATIONS AND FUTURE WORKS

Limitation: Dependency on image backbone. Within our MV-Adapter, we only fine-tune the additional multi-view attention and image cross-attention layers, and do not disturb the original structure or feature space. Consequently, the overall performance of MV-Adapter is heavily dependent on the base T2I model. If the foundational model struggles to generate content that aligns with the provided prompt or produces images of low quality, MV-Adapter is unlikely to compensate for these deficiencies. On the other hand, employing superior image backbones can enhance the synthetic results. We present a comparison of outputs generated using SDXL (Podell et al., 2024) and SD2.1 (Rombach et al., 2022) models in Fig. 17, which confirms this observation, particularly in text-conditioned multi-view generation. We believe that MV-Adapter can be further developed by utilizing advanced T2I models (Team, 2024; Labs, 2024) based on the DiT architecture (Peebles & Xie, 2023), to achieve higher visual quality in the generated images.



Figure 17: Qualitative comparison of our MV-Adapter based on SD2.1 and SDXL.

Future works: 3D scene generation, dynamic multi-view video generation, inspiration for modeling new knowledge. This paper provides extensive analyses and enhancements for our novel multi-view adapter, MV-Adapter. While our model has significantly improved efficiency, adaptability, versatility, and performance compared to previous models, we identify several promising areas for future work:

- 3D scene generation. Our method can be extended to scene-level multi-view generation, accommodating both camera- and geometry-guided approaches with text or image conditions.
- Dynamic multi-view video generation. Exploring dynamic multi-view video generation using a similar approach as MV-Adapter within text-to-video generation models (Zheng et al., 2024; Yang et al., 2024) presents a valuable opportunity for further advancements.
- Inspiration for modeling new knowledge. Our approach of decoupling the learning of geometric knowledge from the image prior can be applied to learning zoom in/out effects, consistent lighting, and other viewpoint-dependent properties. It also provides valuable insights for modeling physical or temporal knowledge based on image priors.

A.5 MORE COMPARISON RESULTS

A.5.1 IMAGE-TO-MULTI-VIEW GENERATION

To provide a more in-depth analysis of our quantitative results on image-to-multi-view generation, we conducted a user study comparing MV-Adapter (based on SD2.1 (Rombach et al., 2022)) with baseline methods (Wang & Shi, 2023; Shi et al., 2023a; Wang et al., 2024b; Voleti et al., 2024; Wen et al., 2024; Li et al., 2024). The study aimed to evaluate both multi-view consistency and image quality preferences. We selected 30 samples covering a diverse range of categories, such as toy cars, medicine bottles, stationery, dolls, and sculptures. A total of 50 participants were recruited to provide their preferences between the outputs of different methods.

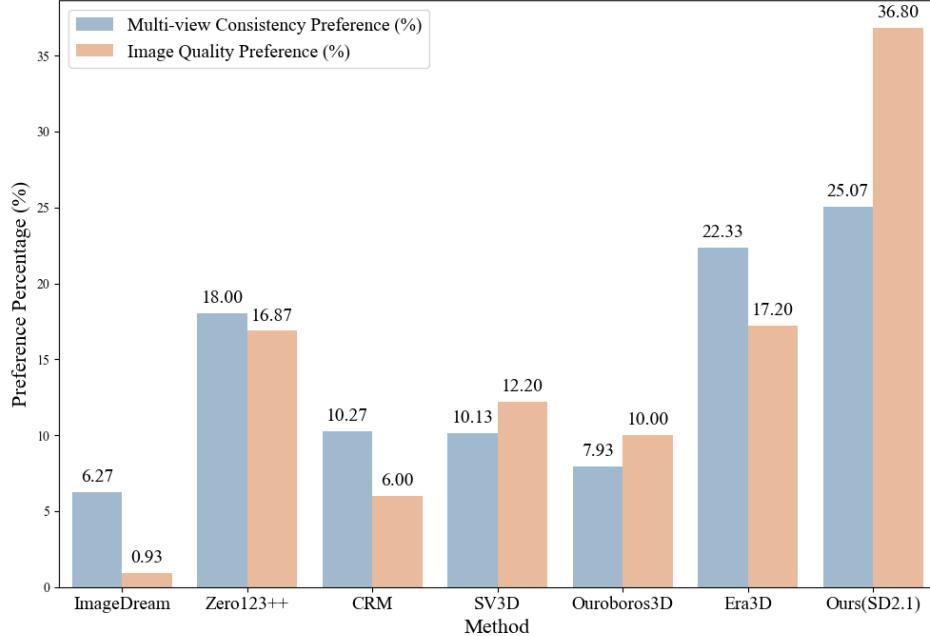


Figure 18: Results of user study on image-to-multi-view generation.

Participants were presented with pairs of multi-view images generated by MV-Adapter and the baseline methods. For each pair, they were asked to choose the one they preferred in terms of multi-view consistency and image quality. The results of the user study are summarized in Fig. 18. The findings indicate that, in terms of multi-view consistency, MV-Adapter performs comparably to Era3D, with preference rates of 25.07% and 22.33%, respectively. However, regarding image quality, MV-Adapter demonstrates a significant advantage, receiving a higher preference rate of 36.80% compared to the baseline methods. The improved image quality can be attributed to MV-Adapter’s ability to leverage the strengths of the underlying T2I models without full fine-tuning, preserving the original feature space and benefiting from the high-quality priors of the base models.

A.6 MORE VISUAL RESULTS

In Fig. 19 and Fig. 20, we show more visual results of MV-Adapter on camera-guided text-to-multiview generation with community models and extensions, such as ControlNet (Zhang et al., 2023) and IP-Adapter (Ye et al., 2023). In Fig. 21, we show more visual results on camera-guided image-to-multiview generation. In Fig. 22, we show more visual results on text-to-3D generation. In Fig. 23, we show more visual results on image-to-3D generation. In Fig. 24, we show more visual results on geometry-guided text-to-texture generation. In Fig. 25, we show more visual results on geometry-guided image-to-texture generation. Note that we have removed the background of the generated images in the visual results.

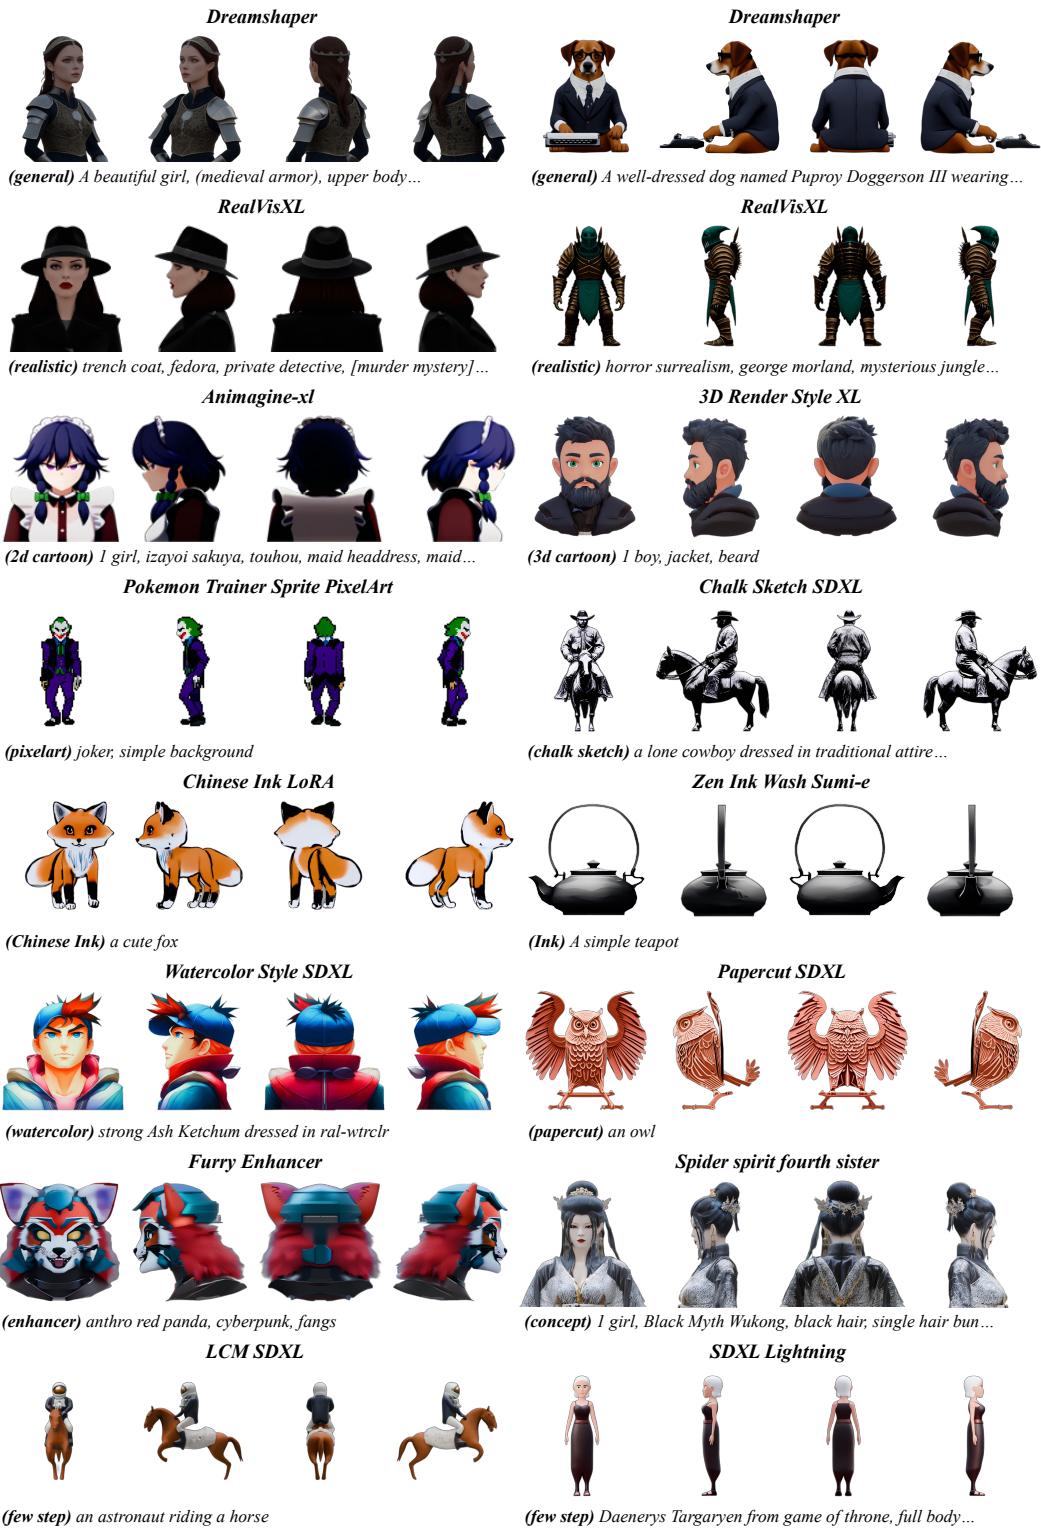


Figure 19: Additional results on camera-guided text-to-multiview generation with community models.

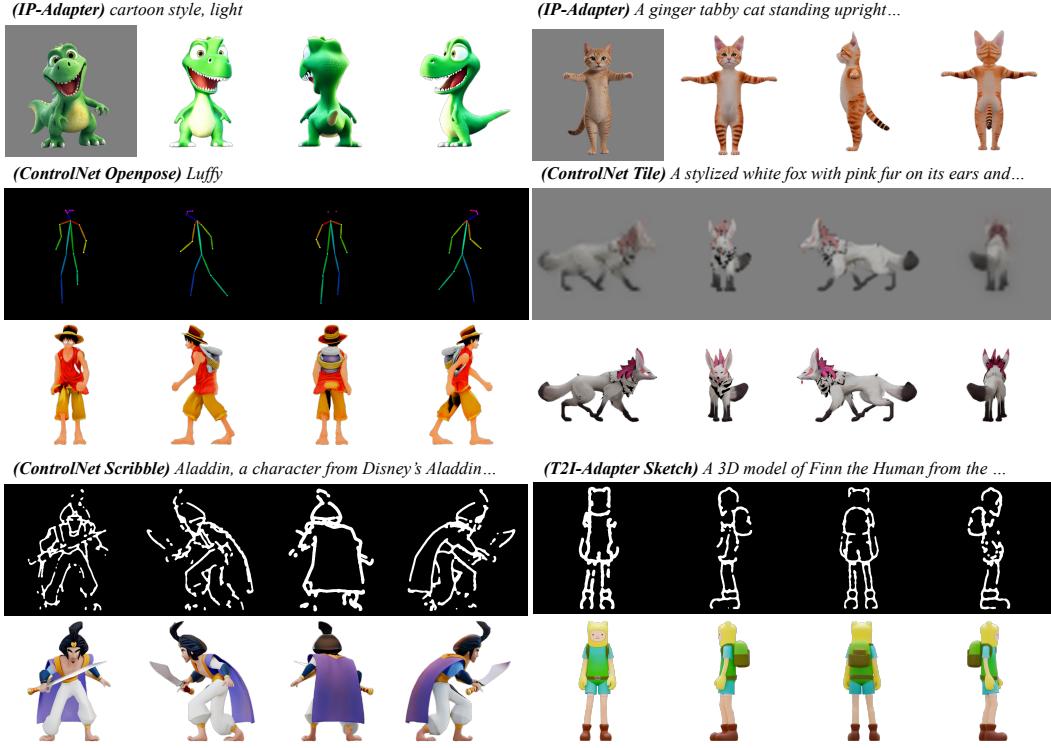


Figure 20: Additional results on camera-guided text-to-multiview generation with extensions.

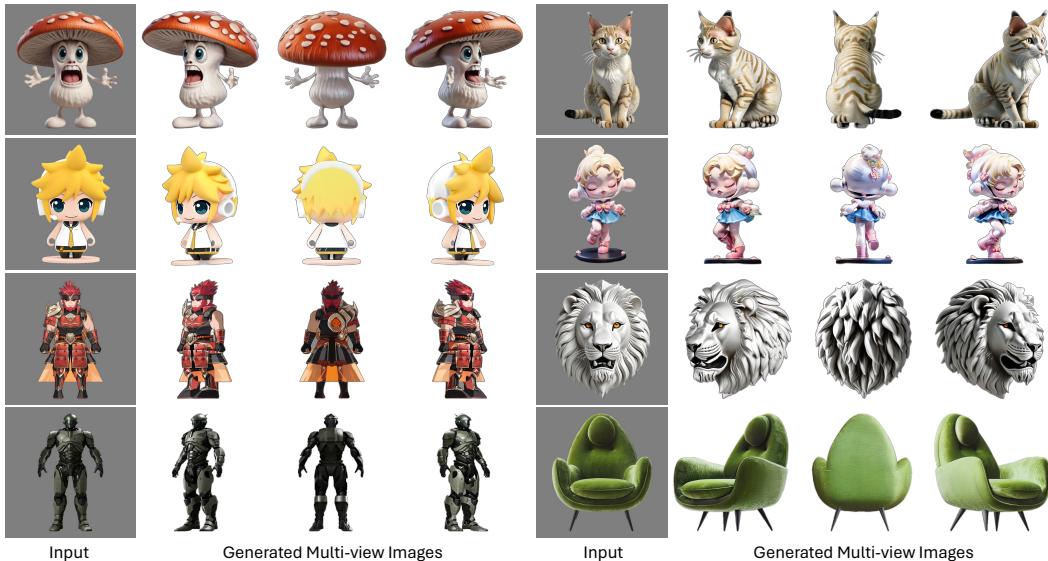


Figure 21: Additional results on camera-guided image-to-multiview generation.



Figure 22: Visual results on text-to-3D generation.

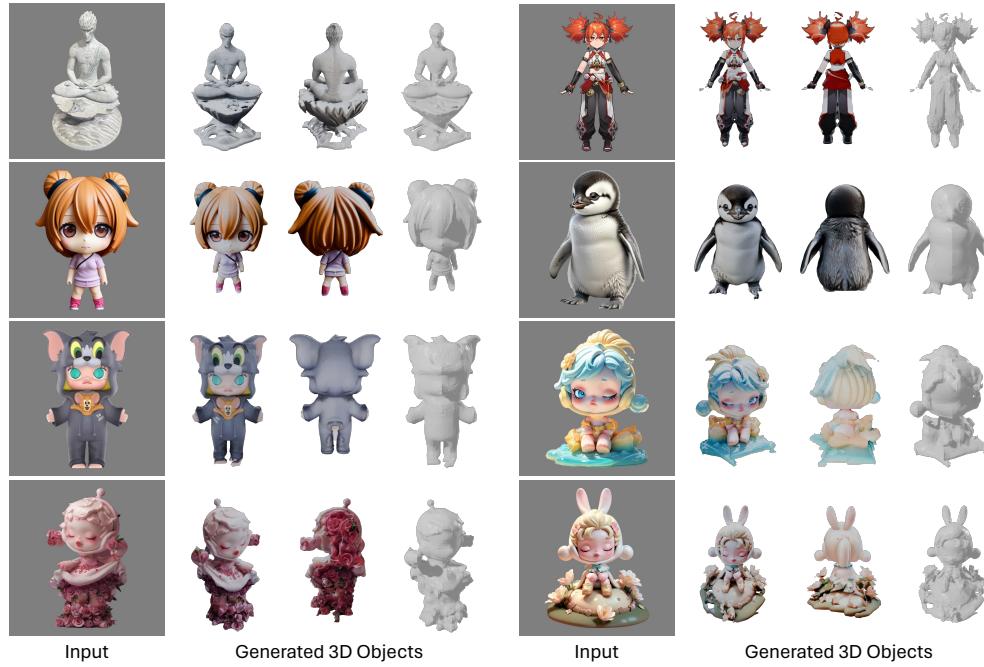


Figure 23: Visual results on image-to-3D generation.

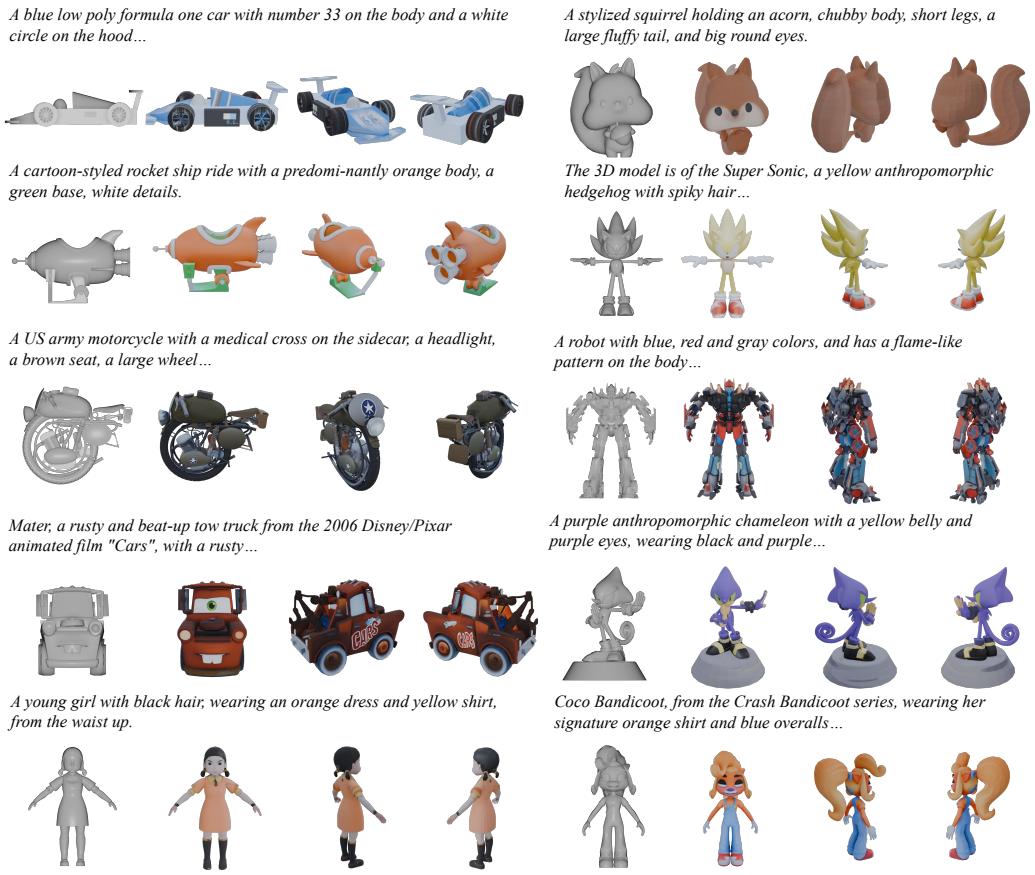


Figure 24: Additional results on geometry-guided text-to-texture generation.

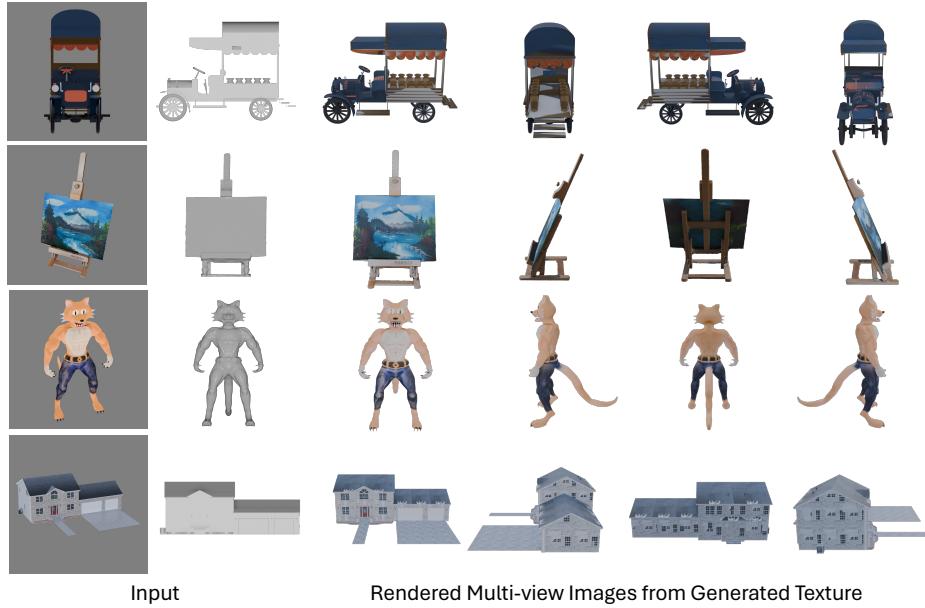


Figure 25: Additional results on geometry-guided image-to-texture generation.